

Chapter 1: Intro

OS is a program that acts as an intermediary between a user of a computer and the computer hardware.
Goals: Execute user programs and make solving user problems easier; computer sys convenient to use; efficiently use hardware; hide/hardware peculiarities. (Res allocator, etc program)
Computer sys: Hardware (CPU, mem, I/O device; basic computing resources), OS (Ctl & coordinate use of hardware among apps and users), apps (computers), user
Computer startup: bootstrap program (EPROM, firmware CPU & I/O: execute concurrently; Each device controller has a local buffer; CPU moves data from/to memory to/from local buffers; I/O is from the device to local buffer of controller; I/O interrupt).
Interrupt: transfer control to ISR (interrupt service routine); mark the addr of the interrupted instruction; interrupt vector contains the addr of all the service routines; OS is interrupt driven; interrupt request (IR)

Trap & Exception: software interrupt. (Intercept is hardware)
Type of interrupts: Polling(CPU check)& vectored interrupt system
Asynchronous I/O: Sys call, Device-status table (IO device's type, address, and state)

Mem-device hierarchy: main mem (RA, volatile), secondary storage (non-volatile), solid state disks, magnetic disk (non-v, tracts, sections)
Caching: Information in use copied from slower to faster storage temporarily. Cache size and replacement policy.

Device Driver: provides uniform interface between controller and kernel
Direct Mem Access: Only one interrupt is generated per block, rather than the one interrupt per byte

Processors: single general purpose processor; multiprocessors/parallel system -> throughput, economy of scale, graceful degradation or fault tolerance. (sync, async)
Clustered Sys: Like multiprocessor systems, but multiple systems working together; storage area network; high-availability service which survives failures (sync, async clustering); HPC; Distributed Lock Manager (conflict ops)
Dual mode: kernel and user mode; mode bit; sys call change to kernel mode and return to user mode; VMM mode.

Process: Program is a passive entity, process is an active entity. A process is a program in execution. One thread, one PCounter
Process management: create&delete user/sys processes; suspending and resuming; p sync, p communicate; deadlock

Memory management: keep track of usage and by whom; move in and out; allocate and deallocate.

Storage Management: abstract physical properties to logical storage unit - file.

I/O subsys: Memory management of I/O including buffering (storing data temporarily while it is being transferred), caching (storing parts of data in fast storage for performance), spooling (the overlapping of output of one job with input of other jobs); General/specific device-driver interface

Protection and Security: user id, group id, privilege escalation
Computing environment: traditional, mobile, distributed, client-server, peer-to-peer, virtualization, cloud computing, real-time embedded sys.

Chapter 2: Sys Structure

OS services: UI (GUI, TUI, CLI, Batch), program exec, IO op, file-sys, communication, error detection, res allocation, accounting, protection and security.

Sys calls: programming interface to services provided by OS; in C/C++; accessed via API rather than direct use (hide details)

Passing parameters: 1) pass to reg 2) store in a block, table in mem, pass the block addr 3) onto the stack by program and pop by OS.

Sys programs: in user's view, OS is defined by system programs, not sys calls; file management; statinfo; file modification; programming-lang-support; program-load and exec; communications; background services; app programs
OS design goals: 1) user convenient to use, easy to learn, reliable, safe, and fast 2) sys:easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient.

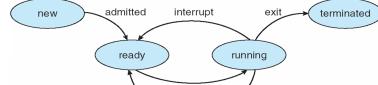
Micro kernel: extend port OS to new architectures; more reliable and secure. Downside: performance overhead
Loadable Kernel Modules: most popular OS implementation
Dynamically loadable kernel: kernel extensions
Monolithic kernel:

iOS: Cocoa Touch, Media Service, Core Service, Core OS Kernighan's Law. "Debugging is twice as hard as writing the code in the first place. Therefore, if you write the code as cleverly as possible, you are, by definition, not smart enough to debug it." (log files; core dump; crash dump; tracing list; profiling tool)

Sys Booting: bootstrap loader in ROM/EROM locates the kernel; two step, boot block in ROM which then load the boot loader (commonly GRUB->allow kernel selection) from disk.

Chapter 3: Process Concept

Process: text section, program counter, stack, data section, heap



Process Control Block (PCB): process state, PC, CPU reg, CPU scheduling info, mem-manage info, accounting info, I/O

Process Scheduling: ready/running/Device queue

Long-term Scheduler: select which process -> ready (frequently invoked - milliseconds) control the degree of multi-programming

Short-term Scheduler: select which process -> exec & CPU (infrequently invoked - seconds, minutes)

CPU-bound / I/O-bound process: longterm sch finds good mix

Medium Term Scheduling: swap in/partially executed swapped-out processes swap out

Context Switch: ready queue -> CPU -> I/O queue -> time slice expired -> child executes -> fork a child -> interrupt occurs -> waiting for an interrupt -> ready queue

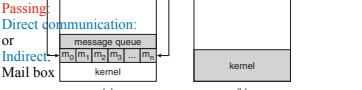
Dynamic relocation: Static Linking: system libraries and program code combined by the loader into the binary program image

Dynamic Linking: linking postponed until execution time; use a piece of code (stab) to locate and replace itself with the appropriate memory-resident library routine. OS check if routine is in the addr space add if not. Particular useful to libs.

Swapping: A process can be swapped temporarily out of memory to a backing store, and then brought back into memory for continued execution. **Backing store:** fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images.

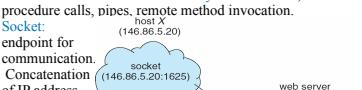
Process Creation: parent creates child which in turn creates child. (Resource Sharing / Execution Options): The address space duplicated and child program is loaded into it. (fork/exec/wait)
Process Identifier: pid used to manage processes.
Process Termination: exit() -> ask the OS to delete it; abort() - parent may terminate the child
Cascading Termination: terminate all children when parent term
Zombie: a terminated child w/o parent waiting
Orphans: if parent terminates, the process is an orphan.
Cooperating Process: supports info sharing, computation speedup, modularity, convenience -> needs interprocess communication (IPC).

IPC models: shared mem(bounded/unb buffer), message passing,



(a) Synchronization: Message passing may be blocking (sync) (sender/receiver blocked until received or msg available) or non-b (async) (sender send and continue, receiver may continue with null)
Buffering: zero capacity (sender must wait receiver), bounded capacity (finite len of msg, sender wait if full), unbounded c.

Communication in Client-Server Systems: sockets, remote procedure calls, pipes, remote method invocation.



Three types of Sockets:

Connection-oriented(TCP), Connectionless(UDP), MulticastSocket class.

Remote Procedure Call: stubs (client side) locate the server and marshal the parameters, server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server

Chapter 4: MultiThread Programming

Motivation: process is heavy weight, simplify code, increase efficiency; kernels are generally multithreaded.

Processes: are completely separated program with its own variables, stacks and mem allocation while **Threads** shares the same mem space and global variables between routines.

Benefits: 1) responsiveness: May allow continued execution if part of process is blocked, especially important for user interfaces 2) **resource sharing** 3) **economy:** Cheaper than process creation, thread switching lower overhead than context switching 4) **scalability:** may use multi-processor architecture

Concurrency: core 1: T1, T3, T1, T3, T1, ...
Parallelism: core 2: T2, T4, T2, T4, T2, ...

Amdahl's Law: speedup <= 1 / (S + (1-S)/N) where S - the serial portion, N - the processing cores.

User thread & Kernel thread: Many to One (Many user-level threads mapped to single kernel thread, one thread blocked causes all to block). One to One (more concurrent). Many to Many (OS creates sufficient number of kernel threads) two level

Thread: may be provided either as user level or kernel level

```
pthread_mutex_lock(&mutex1);
while (c < 0) {
    pthread_cond_wait(&cond1, &mutex1);
    if (c == 0) {
        pthread_cond_signal(&cond1);
    }
    pthread_mutex_unlock(&mutex1);
    action();
}
```

Signals are not remembered

Implicit Threading: creation and management of threads done by compilers and runtime libraries rather than programmers (methods: Thread Pools, OpenMP, GrandCentral Dispatch)

Thread Pools: create a number of threads in a pool where they await work. **Advantages:** 1) slightly faster to service a request with an existing thread than create a new thread 2) # of threads in app bound to the size of the pool 3) separate task from mechanics of create tasks-different strategies for running task

OpenMP: set of compiler derivatives and an API for C, C++, FORTRAN; Support for parallel programming in shared-memory environments; Identifies parallel regions - blocks of code that can run in parallel.

Grand Central Dispatch: Apple technology for Mac OS X and iOS operating systems. Similar to OpenMP's function.

Thread Issues: fork (duplicate all?)/exec; **Signal Handling** (handler routine: deliver to > process; thread which the signal applies every thread; certain thread; assign a thread to take all);

Thread-Cancellation (asynchronous / deferred cancellation); Thread-Local Storage (allows each thread to have its own copy of data, similar to static data); **Scheduler Activities** (appropriate number of kernel threads allocated to the application (1:m:1); use an intermediate data structure between user and kernel threads - lightweight process (LWP))

Chapter 8: Memory Management

OS: Process/Mem/Storage/File/I/O management

Background: Program must be brought (from disk) into memory and placed within a process for it to run; Main memory and registers are only storage CPU can access directly.

Base and Limit Reg: define the logical address space. CPU checks every mem access from user mode to make sure it is between the base and limit for that user.

Addr Binding: Compile time (addr known at compile time - absolute code). Load time (must generate relocatable code), Execution time (allow process to be moved during its exec)

Multi-Step Processing of a user program: src -> [compiler/assembler] -> object module -> [linkage editor] (other object modules comes in) -> load module -> [loader] (sys libs comes in) -> [in-mem binary mem image] (dynamic loaded sys lib)

Logical & Physical addr: same in compile-time and load time addr-binding schema. Different in exec-time addr binding.

User program deals with logical addr, it never sees physical addr. Execution-time binding occurs when reference is made to location in memory.

Mem management Unit: hardware device that at run time maps virtual to physical addr.

Dynamic relocation.

Static Linking: system libraries and program code combined by the loader into the binary program image

Dynamic Linking: linking postponed until execution time; use a piece of code (stab) to locate and replace itself with the appropriate memory-resident library routine. OS checks if routine is in the address space add if not. Particular useful to libs.

Swapping: A process can be swapped temporarily out of memory to a backing store, and then brought back into memory for continued execution. **Backing store:** fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images.

Roll out Roll in: swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed. Depending on the binding method and IO pending, the process may be swapped back to the exact same addr or not.

Overhead: proportional to the swap size.

Modern OS: Swap only when free memory extremely low

Swapping on Mobile sys: Not typically supported, but has flash for quick restart.

Contiguous Allocation: OS usually held in low mem with interrupt vector. User process in high addr. Each process contained in a single contiguous mem section. Relocation registers (base, limit, relocation) used to protect user processes from each other, and from changing operating-system code and data.

Multiple Partition Allocation: limit the degree of multiprogramming; Variable-partition size; Hole (fragmentation)

First fit / Best fit / Worst fit:

External Fragmentation: total memory space exists to satisfy a request, but it is not contiguous (and none of them is large enough to fit the program alone) -> **compaction:** Shuffle memory contents to place all free memory together in one large block.

Internal Fragmentation: size difference between allocated mem and used mem internal to a partition, but not being used.

Segmentation: logical addr <segment num, offset>; very length

Segmentation table: base contains the starting physical address where the segments reside in memory; limit specifies the length of the segment. Each entry is associated with validation

Segmentation base/length reg (STBR, STLR):

Paging: PM -> frames; LM -> pages; Page table to translate; N pages need N free frames to be loaded. **Backing store** likewise split into pages. Still has **fragmentation**.

Page Num & Page Offset:

Trade-off: small page size, less fragmentation but larger page table

Page Table Implementation: Page table is kept in the mem.

Page table base / length reg: this scheme every data/instruction access requires two memory accesses. **Translation look-aside buffers / Associative Reg (TLBs)** are used for the two mem access procedure. **Address-space identifiers (ASIDs)** may be stored in each entry.

1. First visit on mem is on page table,取出虚页对应的物理页。

2. 第二次访问内存是访问实际内存地址。1. 第一次访问TLB, 得到虚页对应的物理页

3. 第二次访问的是内存, 访问实际地址。

4. 第一次访问TLB, 得到虚页对应的物理页

5. 第二次访问的是内存, 访问实际地址。

6. 第一次访问TLB, 得到虚页对应的物理页

7. 第二次访问的是内存, 访问实际地址。

8. 第一次访问TLB, 得到虚页对应的物理页

9. 第二次访问的是内存, 访问实际地址。

10. 第一次访问TLB, 得到虚页对应的物理页

11. 第二次访问的是内存, 访问实际地址。

12. 第一次访问TLB, 得到虚页对应的物理页

13. 第二次访问的是内存, 访问实际地址。

14. 第一次访问TLB, 得到虚页对应的物理页

15. 第二次访问的是内存, 访问实际地址。

16. 第一次访问TLB, 得到虚页对应的物理页

17. 第二次访问的是内存, 访问实际地址。

18. 第一次访问TLB, 得到虚页对应的物理页

19. 第二次访问的是内存, 访问实际地址。

20. 第一次访问TLB, 得到虚页对应的物理页

21. 第二次访问的是内存, 访问实际地址。

22. 第一次访问TLB, 得到虚页对应的物理页

23. 第二次访问的是内存, 访问实际地址。

24. 第一次访问TLB, 得到虚页对应的物理页

25. 第二次访问的是内存, 访问实际地址。

26. 第一次访问TLB, 得到虚页对应的物理页

27. 第二次访问的是内存, 访问实际地址。

28. 第一次访问TLB, 得到虚页对应的物理页

29. 第二次访问的是内存, 访问实际地址。

30. 第一次访问TLB, 得到虚页对应的物理页

31. 第二次访问的是内存, 访问实际地址。

32. 第一次访问TLB, 得到虚页对应的物理页

33. 第二次访问的是内存, 访问实际地址。

34. 第一次访问TLB, 得到虚页对应的物理页

35. 第二次访问的是内存, 访问实际地址。

36. 第一次访问TLB, 得到虚页对应的物理页

37. 第二次访问的是内存, 访问实际地址。

38. 第一次访问TLB, 得到虚页对应的物理页

39. 第二次访问的是内存, 访问实际地址。

40. 第一次访问TLB, 得到虚页对应的物理页

41. 第二次访问的是内存, 访问实际地址。

42. 第一次访问TLB, 得到虚页对应的物理页

43. 第二次访问的是内存, 访问实际地址。

44. 第一次访问TLB, 得到虚页对应的物理页

45. 第二次访问的是内存, 访问实际地址。

46. 第一次访问TLB, 得到虚页对应的物理页

47. 第二次访问的是内存, 访问实际地址。

48. 第一次访问TLB, 得到虚页对应的物理页

49. 第二次访问的是内存, 访问实际地址。

50. 第一次访问TLB, 得到虚页对应的物理页

51. 第二次访问的是内存, 访问实际地址。

52. 第一次访问TLB, 得到虚页对应的物理页

53. 第二次访问的是内存, 访问实际地址。

54. 第一次访问TLB, 得到虚页对应的物理页

55. 第二次访问的是内存, 访问实际地址。

56. 第一次访问TLB, 得到虚页对应的物理页

57. 第二次访问的是内存, 访问实际地址。

58. 第一次访问TLB, 得到虚页对应的物理页

59. 第二次访问的是内存, 访问实际地址。

60. 第一次访问TLB, 得到虚页对应的物理页

61. 第二次访问的是内存, 访问实际地址。

62. 第一次访问TLB, 得到虚页对应的物理页

63. 第二次访问的是内存, 访问实际地址。

64. 第一次访问TLB, 得到虚页对应的物理页

65. 第二次访问的是内存, 访问实际地址。

66. 第一次访问TLB, 得到虚页对应的物理页

67. 第二次访问的是内存, 访问实际地址。

68. 第一次访问TLB, 得到虚页对应的物理页

69. 第二次访问的是内存, 访问实际地址。

70. 第一次访问TLB, 得到虚页对应的物理页

71. 第二次访问的是内存, 访问实际地址。

72. 第一次访问TLB, 得到虚页对应的物理页

73. 第二次访问的是内存, 访问实际地址。

74. 第一次访问TLB, 得到虚页对应的物理页

75. 第二次访问的是内存, 访问实际地址。

76. 第一次访问TLB, 得到虚页对应的物理页

77. 第二次访问的是内存, 访问实际地址。

78. 第一次访问TLB, 得到虚页对应的物理页

79. 第二次访问的是内存, 访问实际地址。

80. 第一次访问TLB, 得到虚页对应的物理页

81. 第二次访问的是内存, 访问实际地址。

82. 第一次访问TLB, 得到虚页对应的物理页

83. 第二次访问的是内存, 访问实际地址。

84. 第一次访问TLB, 得到虚页对应的物理页

85. 第二次访问的是内存, 访问实际地址。

86. 第一次访问TLB, 得到虚页对应的物理页

87. 第二次访问的是内存, 访问实际地址。

88. 第一次访问TLB, 得到虚页对应的物理页

89. 第二次访问的是内存, 访问实际地址。

90. 第一次访问TLB, 得到虚页对应的物理页

91. 第二次访问的是内存, 访问实际地址。

